

# Credit Card Fraud Detection

## Importing necessary libraries

```
In [46]: import numpy as np
import pandas as pd

from sklearn.model_selection import train_test_split
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.utils import resample
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import classification_report, accuracy_score, confusion_matrix
from sklearn.model_selection import cross_val_score
```


## Importing the dataset

```
In [2]: df_creditcard = pd.read_csv("C:/Users/c2108436/OneDrive - Teesside University/credit-card-fraud-Copy1.csv")
df_creditcard.head()
```

```
Out[2]:
```

	Unnamed: 0	trans_date_trans_time	cc_num	merchant	category	amt
0	0	2020-06-21 12:14:25	2291163933867244	fraud_Kirlin and Sons	personal_care	2.86
1	1	2020-06-21 12:14:33	3573030041201292	fraud_Sporer-Keebler	personal_care	29.84
2	2	2020-06-21 12:14:53	3598215285024754	fraud_Swaniawski, Nitzsche and Welch	health_fitness	41.28
3	3	2020-06-21 12:15:15	3591919803438423	fraud_Haley Group	misc_pos	60.05
4	4	2020-06-21 12:15:17	3526826139003047	fraud_Johnston-Casper	travel	3.19

5 rows × 23 columns



## Summary of the dataframe

In [3]: `df_creditcard.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 555719 entries, 0 to 555718
Data columns (total 23 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Unnamed: 0                            555719 non-null  int64
1   trans_date_trans_time                 555719 non-null  object
2   cc_num                               555719 non-null  int64
3   merchant                             555719 non-null  object
4   category                             555719 non-null  object
5   amt                                   555719 non-null  float64
6   first                                555719 non-null  object
7   last                                  555719 non-null  object
8   gender                               555719 non-null  object
9   street                               555719 non-null  object
10  city                                  555719 non-null  object
11  state                                555719 non-null  object
12  zip                                   555719 non-null  int64
13  lat                                   555719 non-null  float64
14  long                                  555719 non-null  float64
15  city_pop                             555719 non-null  int64
16  job                                   555719 non-null  object
17  dob                                   555719 non-null  object
18  trans_num                            555719 non-null  object
19  unix_time                            555719 non-null  int64
20  merch_lat                            555719 non-null  float64
21  merch_long                           555719 non-null  float64
22  is_fraud                             555719 non-null  int64
dtypes: float64(5), int64(6), object(12)
memory usage: 97.5+ MB
```

## Data Exploration/cleaning

In [4]: `df_creditcard.shape` *# To check the dimension*

Out[4]: (555719, 23)

```
In [5]: # statistical summary
df_creditcard.describe()
```

```
Out[5]:
```

	Unnamed: 0	cc_num	amt	zip	lat	lon
count	555719.000000	5.557190e+05	555719.000000	555719.000000	555719.000000	555719.000000
mean	277859.000000	4.178387e+17	69.392810	48842.628015	38.543253	-90.231000
std	160422.401459	1.309837e+18	156.745941	26855.283328	5.061336	13.721000
min	0.000000	6.041621e+10	1.000000	1257.000000	20.027100	-165.672000
25%	138929.500000	1.800429e+14	9.630000	26292.000000	34.668900	-96.798000
50%	277859.000000	3.521417e+15	47.290000	48174.000000	39.371600	-87.476000
75%	416788.500000	4.635331e+15	83.010000	72011.000000	41.894800	-80.175000
max	555718.000000	4.992346e+18	22768.110000	99921.000000	65.689900	-67.950000

## To view the data type of each features

```
In [6]: df_creditcard.dtypes
```

```
Out[6]: Unnamed: 0          int64
trans_date_trans_time    object
cc_num                   int64
merchant                 object
category                 object
amt                      float64
first                    object
last                     object
gender                   object
street                   object
city                     object
state                    object
zip                      int64
lat                      float64
long                     float64
city_pop                 int64
job                      object
dob                      object
trans_num                object
unix_time                int64
merch_lat                float64
merch_long               float64
is_fraud                 int64
dtype: object
```

## Checking if any missing values present in the dataset

```
In [7]: df_creditcard.isnull().sum()
```

```
Out[7]: Unnamed: 0      0
trans_date_trans_time  0
cc_num                 0
merchant               0
category               0
amt                   0
first                 0
last                  0
gender                0
street                0
city                  0
state                 0
zip                   0
lat                   0
long                  0
city_pop              0
job                   0
dob                   0
trans_num             0
unix_time             0
merch_lat             0
merch_long            0
is_fraud              0
dtype: int64
```

```
In [8]: df_creditcard.duplicated().sum() # check for duplicates
```

```
Out[8]: 0
```

## Determine class distribution of the target variable

```
In [9]: y = df_creditcard['is_fraud']


print(f'Percentage of fraudulent transactions: % {round(y.value_counts(normalized=True).get(1) * 100, 2)}')
```

```
Percentage of fraudulent transactions: % 0.39 --> (2145 transactions)
Percentage of genuine transactions: % 99.61 --> (553574 transactions)
```

## We have an unbalanced data


```
In [10]: import seaborn as sns  
  
sns.countplot(x= df_creditcard["is_fraud"])
```

```
Out[10]: <Axes: xlabel='is_fraud', ylabel='count'>
```



```
In [11]: class_counts = df_creditcard['is_fraud'].value_counts()
plt.pie(class_counts, labels=class_counts.index, autopct='%1.1f%%')
plt.title('Class Distribution showing imbalanced Dataset')
plt.show()
```

Class Distribution showing imbalanced Dataset



```
In [12]: # Separate the fraudulent transactions from the genuine ones
df_creditcard_majority = df_creditcard[df_creditcard.is_fraud==0]
df_creditcard_minority = df_creditcard[df_creditcard.is_fraud==1]
```


**Let's resample the minority class to balance the dataset using `resample()`**

```
In [13]: df_creditcard_minority_upsampled = resample(df_creditcard_minority, replace=
                                                    n_samples=len(df_creditcard_majority),
                                                    random_state=42)    # reproducible results
```

```
In [14]: #Let's combine the upsampled minority class with the majority class
df_creditcard_balanced = pd.concat([df_creditcard_majority, df_creditcard_mi
```


```
In [15]: #lets verify and see if the dataset is now balanced
sns.countplot(x= df_creditcard_balanced["is_fraud"])
```

Out[15]: <Axes: xlabel='is\_fraud', ylabel='count'>

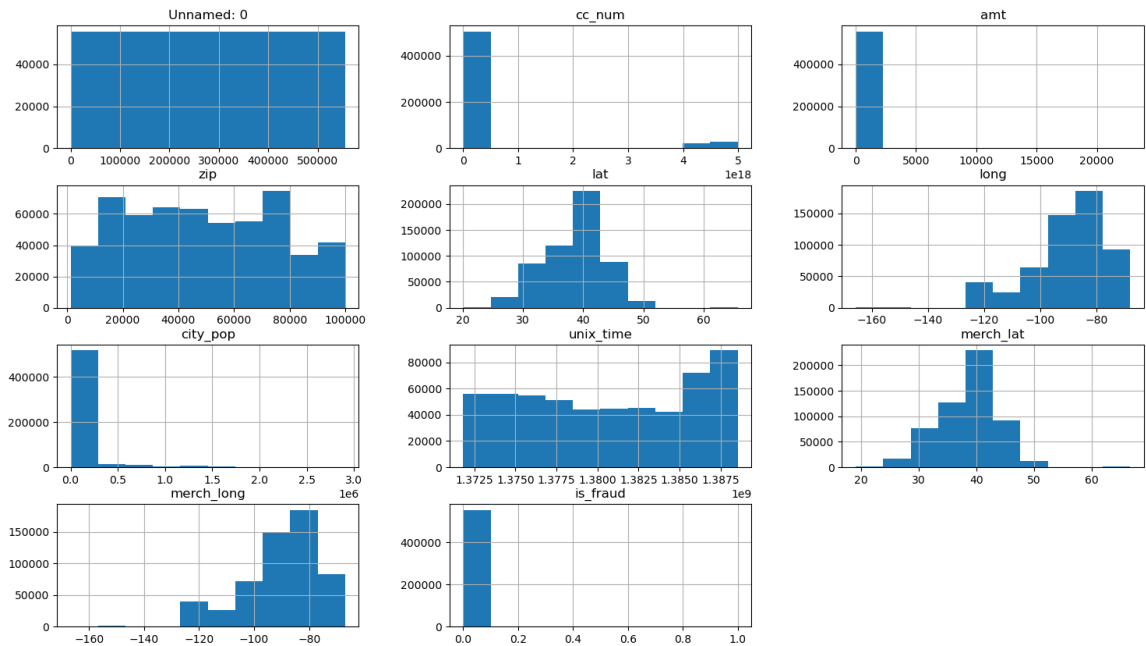


```
In [16]: # Lets visualize the balanced dataset using a pie chart
class_counts_balanced = df_creditcard_balanced['is_fraud'].value_counts()
plt.pie(class_counts_balanced, labels=class_counts_balanced.index, autopct='
plt.title('Class Distribution of the Balanced Dataset')
plt.show()
```

Class Distribution of the Balanced Dataset



```
In [17]: df_creditcard.hist(figsize=(18, 10));
# Histogram to show the skewness in the dataset
```



## Feature engineering

The columns like 'cc\_num', 'first', 'last', 'trans\_num' don't provide significant relevant information related to fraud detection. So, we drop it.

```
In [18]: print(df_creditcard.columns)
df_creditcard.drop(['Unnamed: 0', 'first', 'last', 'trans_num', 'job'], axis = 1,
print(df_creditcard.columns)
```

```
Index(['Unnamed: 0', 'trans_date_trans_time', 'cc_num', 'merchant', 'category',
      'amt', 'first', 'last', 'gender', 'street', 'city', 'state', 'zip',
      'lat', 'long', 'city_pop', 'job', 'dob', 'trans_num', 'unix_time',
      'merch_lat', 'merch_long', 'is_fraud'],
      dtype='object')
Index(['trans_date_trans_time', 'cc_num', 'merchant', 'category', 'amt',
      'gender', 'street', 'city', 'state', 'zip', 'lat', 'long', 'city_pop',
      'dob', 'unix_time', 'merch_lat', 'merch_long', 'is_fraud'],
      dtype='object')
```

**Let's convert the transaction date and time to separate columns which includes hour, day of a week and month**

```
In [19]: df_creditcard['trans_date_trans_time'] = pd.to_datetime(df_creditcard['trans_date_trans_time'])
df_creditcard['Hour'] = df_creditcard['trans_date_trans_time'].dt.hour
df_creditcard['Day_of_week'] = df_creditcard['trans_date_trans_time'].dt.dayofweek
df_creditcard['Month'] = df_creditcard['trans_date_trans_time'].dt.month
```



```
In [20]: df_creditcard['dob']=pd.to_datetime(df_creditcard['dob'])
df_creditcard['dob']
```

```
Out[20]: 0      1968-03-19
1      1990-01-17
2      1970-10-21
3      1987-07-25
4      1955-07-06
...
555714 1966-02-13
555715 1999-12-27
555716 1981-11-29
555717 1965-12-15
555718 1993-05-10
Name: dob, Length: 555719, dtype: datetime64[ns]
```

## Let's look at the Frequency of Transactions

```
In [21]: def last1dayTransCnt(df_creditcard):
temp = pd.Series(df_creditcard.index,index=df_creditcard.trans_date_tran
#data (parameter) is df_creditcard.index
#temp is a series whose index is time stamp and value is row indices of
In_a_Day = temp.rolling('1d').count()-1
#in a day is a series with timestamp as index and frequency as its value
In_a_Day.index= temp.values
#in a day 's index is just 0 1 2; row indices of df_creditcard or x
df_creditcard['In_a_Day'] = In_a_Day.reindex(df_creditcard.index)
#df_creditcard
return df_creditcard
```

```
In [22]: def last1weekTransCnt(x):
temp = pd.Series(x.index,index=x.trans_date_trans_time,name="In_a_Week")
In_a_Week = temp.rolling('7d').count()-1
In_a_Week.index = temp.values
x['In_a_Week'] = In_a_Week.reindex(x.index)
return x
```

```
In [23]: def last1monthTransCnt(x):
temp = pd.Series(x.index,index=x.trans_date_trans_time,name="In_a_Month")
In_a_Month = temp.rolling('30d').count()-1
In_a_Month.index = temp.values
x['In_a_Month'] = In_a_Month.reindex(x.index)
return x
```

```
In [24]: df1d = df_creditcard.groupby('cc_num').apply(last1dayTransCnt)
#drop = true ; we don't want to add new column
df1w = df1d.reset_index(drop=True).groupby('cc_num').apply(last1weekTransCnt)
df1dm = df1w.reset_index(drop=True).groupby('cc_num').apply(last1monthTransCnt)
df1dm
```

C:\Users\c2108436\AppData\Local\Temp\ipykernel\_5912\3752136529.py:1: FutureWarning: Not prepending group keys to the result index of transform-like apply. In the future, the group keys will be included in the index, regardless of whether the applied function returns a like-indexed object.  
To preserve the previous behavior, use

```
>>> .groupby(..., group_keys=False)
```

To adopt the future behavior and silence this warning, use

```
>>> .groupby(..., group_keys=True)
df1d = df_creditcard.groupby('cc_num').apply(last1dayTransCnt)
C:\Users\c2108436\AppData\Local\Temp\ipykernel_5912\3752136529.py:3: FutureWarning: Not prepending group keys to the result index of transform-like apply. In the future, the group keys will be included in the index, regardless of whether the applied function returns a like-indexed object.  
To preserve the previous behavior, use
```

```
>>> .groupby(..., group_keys=False)
```

To adopt the future behavior and silence this warning, use

```
>>> .groupby(..., group_keys=True)
df1w = df1d.reset_index(drop=True).groupby('cc_num').apply(last1weekTransCnt)
C:\Users\c2108436\AppData\Local\Temp\ipykernel_5912\3752136529.py:4: FutureWarning: Not prepending group keys to the result index of transform-like apply. In the future, the group keys will be included in the index, regardless of whether the applied function returns a like-indexed object.  
To preserve the previous behavior, use
```

```
>>> .groupby(..., group_keys=False)
```


To adopt the future behavior and silence this warning, use

```
>>> .groupby(..., group_keys=True)
df1dm = df1w.reset_index(drop=True).groupby('cc_num').apply(last1monthTransCnt)
```

Out[24]:

	trans_date_trans_time	cc_num	merchant	category	amt	ger
0	2020-06-21 12:14:25	2291163933867244	fraud_Kirlin and Sons	personal_care	2.86	
1	2020-06-21 12:14:33	3573030041201292	fraud_Sporer-Keebler	personal_care	29.84	
2	2020-06-21 12:14:53	3598215285024754	fraud_Swaniawski, Nitzsche and Welch	health_fitness	41.28	
3	2020-06-21 12:15:15	3591919803438423	fraud_Haley Group	misc_pos	60.05	
4	2020-06-21 12:15:17	3526826139003047	fraud_Johnston-Casper	travel	3.19	
...	...	...	...	...	...	...
555714	2020-12-31 23:59:07	30560609640617	fraud_Reilly and Sons	health_fitness	43.77	
555715	2020-12-31 23:59:09	3556613125071656	fraud_Hoppe-Parisian	kids_pets	111.84	
555716	2020-12-31 23:59:15	6011724471098086	fraud_Rau-Robel	kids_pets	86.88	
555717	2020-12-31 23:59:24	4079773899158	fraud_Breitenberg LLC	travel	7.99	
555718	2020-12-31 23:59:34	4170689372027579	fraud_Dare-Marvin	entertainment	38.13	

555719 rows × 24 columns



As per the frequency of transactions lets guess the fraudulent transactions


## Threshold for a certain transaction to be fraudulent is estimated if no. of day, week or month is more that 90% of the data.

```
In [25]: threshold_day = df1dm['In_a_Day'].quantile(0.9)
threshold_week = df1dm['In_a_Week'].quantile(0.9)
threshold_month = df1dm['In_a_Month'].quantile(0.9)
df1dm['prolly_fraud'] = ((df1dm['In_a_Day']>threshold_day)|
                        (df1dm['In_a_Week']>threshold_week)|
                        (df1dm['In_a_Month']>threshold_month))
df1dm[(df1dm['prolly_fraud'] == True) & (df1dm['is_fraud'] == True)]
```

```
Out[25]:
```

	trans_date_trans_time	cc_num	merchant	category	amt	genc
11799	2020-06-24 23:24:22	180098888332620	fraud_Pfeffer and Sons	shopping_pos	1047.30	
26607	2020-06-30 03:16:01	3586008444788268	fraud_Wilkinson Ltd	entertainment	483.28	
26696	2020-06-30 03:59:50	3586008444788268	fraud_Effertz, Welch and Schowalter	entertainment	520.02	
28628	2020-06-30 16:42:15	3586008444788268	fraud_Schaefer Ltd	kids_pets	19.68	
30010	2020-06-30 23:46:30	3528231451607350	fraud_Torphy-Goyette	shopping_pos	727.32	
...	...	...	...	...	...	...
505774	2020-12-21 02:21:41	4716561796955522	fraud_Murray-Smitham	grocery_pos	358.24	
505826	2020-12-21 02:36:03	4716561796955522	fraud_Schmidt and Sons	shopping_net	859.12	
511244	2020-12-21 22:38:38	4716561796955522	fraud_Quitzon-Goyette	home	209.84	
511272	2020-12-21 22:42:11	4716561796955522	fraud_Schulist Ltd	food_dining	123.58	
511374	2020-12-21 22:59:22	4716561796955522	fraud_Botsford and Sons	home	219.11	

184 rows × 25 columns



## Merged this dataframe's columns to original one

```
In [26]: """df1dm has cc_num as index but we're trying to merge on basis of cc_num cc
so we make cc_num a regular column before merging it"""
df1dm.reset_index(drop =True,inplace=True)
df_creditcard = df_creditcard.merge(df1dm[['trans_date_trans_time','prolly_f
# df.drop(['prolly_fraud_x','prolly_fraud_y'], axis=1, inplace=True)
# df_creditcard.head(3)
```

```
In [27]: df_creditcard.head()
```

```
Out[27]:
```

	trans_date_trans_time	cc_num	merchant	category	amt	gender
0	2020-06-21 12:14:25	2291163933867244	fraud_Kirlin and Sons	personal_care	2.86	M
1	2020-06-21 12:14:33	3573030041201292	fraud_Sporer-Keebler	personal_care	29.84	F
2	2020-06-21 12:14:53	3598215285024754	fraud_Swaniawski, Nitzsche and Welch	health_fitness	41.28	F \
3	2020-06-21 12:15:15	3591919803438423	fraud_Haley Group	misc_pos	60.05	M
4	2020-06-21 12:15:17	3526826139003047	fraud_Johnston-Casper	travel	3.19	M

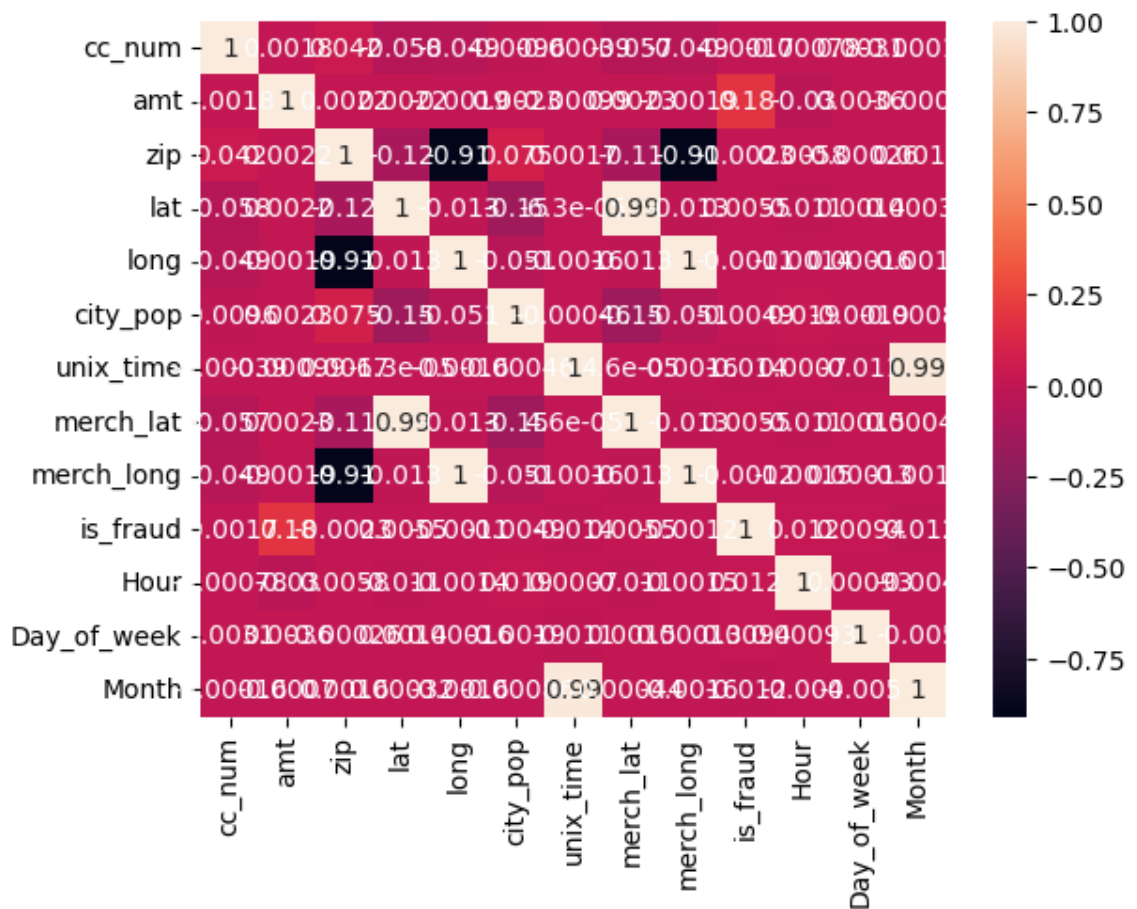
5 rows × 22 columns

## Correlation Heatmaps

### Correlation heatmap of original data

```
In [28]: num_cols = df_creditcard.select_dtypes(include = ['float64','int64'])
sns.heatmap(num_cols.corr(),annot = True)
```


Out[28]: <Axes: >



Correlation heatmap of data with added features

```
In [29]: numeric_columns = df1dm.select_dtypes(include=['float64', 'int64'])
corr_matrix = numeric_columns.corr()
sns.heatmap(corr_matrix, annot=True, fmt='.3f', cmap="YlGnBu")
```

Out[29]: <Axes: >



## Frauds and Categories

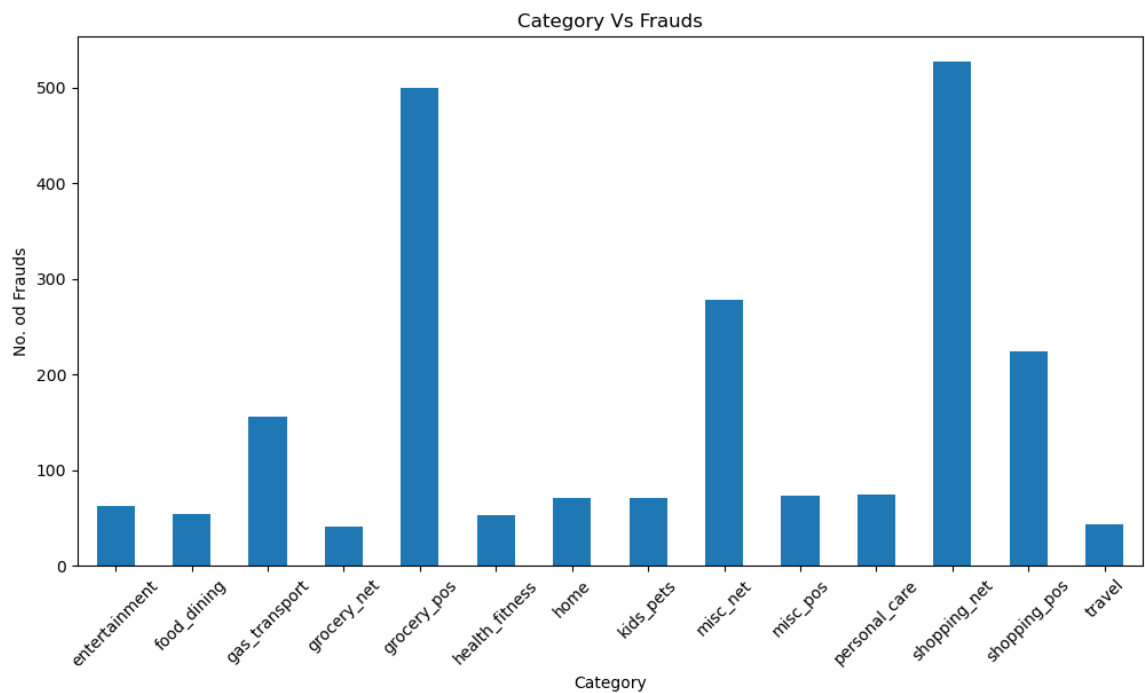
```
In [30]: """
split-apply-combine
first it groups category and is_fraud then it selects is_fraud and counts it
results in
      is_fraud
cat1    0      2
       1      1
cat2    0      1
       1      1
unstacked so the 'is_fraud' gets converted into a single column
is_fraud    0    1
category
cat1         2    1
cat2         1    1
"""
cat_counts = df_creditcard.groupby(['category', 'is_fraud'])['is_fraud'].count()
cat_counts
```

```
Out[30]:
```

	is_fraud	0	1
category			
entertainment		41688	63
food_dining		41055	54
gas_transport		57894	156
grocery_net		19953	41
grocery_pos		53645	500
health_fitness		38377	53
home		54701	71
kids_pets		50896	71
misc_net		27983	278
misc_pos		35769	73
personal_care		41125	74
shopping_net		42933	527
shopping_pos		51585	224
travel		18211	43



```
In [31]: cat_counts_fraud = cat_counts[1]
ccc = cat_counts_fraud.plot(kind='bar', figsize=(12, 6))
ccc.set_ylabel('No. of Frauds')
ccc.set_xlabel('Category')
ccc.set_title('Category Vs Frauds')
plt.xticks(rotation=45)
plt.show()
```

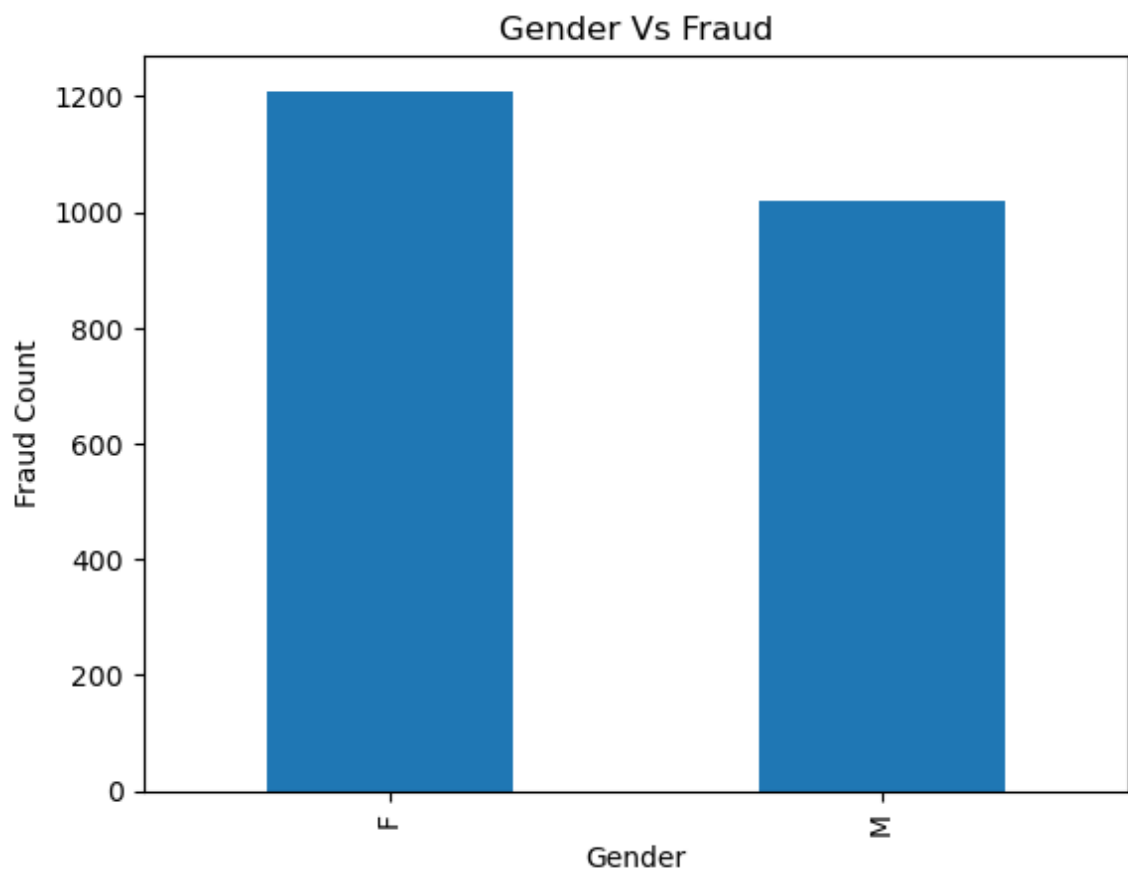


**Grocery and Shopping are the categories with the most frauds.**

## Fraud by Gender

```
In [32]: gen_counts=df_creditcard.groupby(['gender','is_fraud'])['is_fraud'].count()
gen_counts = df_creditcard[df_creditcard['is_fraud'] == 1].groupby('gender')
ax = gen_counts.plot(kind='bar')
ax.set_xlabel('Gender')
ax.set_ylabel('Fraud Count')
ax.set_title('Gender Vs Fraud')
gen_counts
```


```
Out[32]: gender
F      1209
M      1019
Name: is_fraud, dtype: int64
```



The bar chart above shows that females are more involved in credit card fraud

## Most common day of the week for fraud

```
In [33]: fraud_by_day = df_creditcard[df_creditcard['is_fraud']==1].groupby('Day_of_week')
day_labels = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday']
fraud_by_day.index=(day_labels)
fraud_by_day.plot(kind='bar')
plt.ylabel('Number of Fraud Cases')
plt.title('Fraud Cases by Day of the Week')
plt.xticks(rotation = 45)
plt.show()
```



**Most fraud transactions occur during the weekend, especially on Sunday**

## Zip codes based fraud frequency

```
In [34]: df_zip = df_creditcard[df_creditcard['is_fraud'] == 1].groupby('zip')['is_fr  
top_10_zip= df_zip.sort_values(ascending=False).head(10)  
top_10_zip
```

```
Out[34]: zip  
67020      19  
16114      18  
29819      17  
12037      17  
58275      16  
69165      16  
19007      16  
61454      16  
6365       16  
29127      16  
Name: is_fraud, dtype: int64
```

## City Vs Fraud

```
In [35]: df_city = df_creditcard[df_creditcard['is_fraud'] == 1].groupby('city')['is_  
top_10_city= df_city.sort_values(ascending=False).head(10)  
top_10_city
```

```
Out[35]: city  
Camden      29  
Birmingham 26  
Burrton     19  
Clarks Mills 18  
Chatham     17  
Bradley     17  
Preston     16  
Heislerville 16  
Bristol     16  
Jay         16  
Name: is_fraud, dtype: int64
```

**Camden has the highest number of frauds occurrence.**

## Merchants Vs Fraud

```
In [36]: df_mer = df_creditcard[df_creditcard['is_fraud'] == 1].groupby('merchant')['is_fraud'].agg('count').reset_index()
top_10_mer = df_mer.sort_values(ascending=False).head(10)
top_10_mer
```

```
Out[36]: merchant
fraud_Romaguera, Cruickshank and Greenholt    19
fraud_Lemke-Gutmann                           19
fraud_Mosciski, Ziemann and Farrell            19
fraud_Medhurst PLC                           18
fraud_Schultz, Simonis and Little             17
fraud_Heathcote, Yost and Kertzmann           17
fraud_Miller-Hauck                            16
fraud_Heathcote LLC                           16
fraud_Kilback LLC                             15
fraud_Wolf Inc                                15
Name: is_fraud, dtype: int64
```

## Dropping some more features, and encoding categorical features and scaling with numeric values

```
In [37]: X = df_creditcard.drop(['zip', 'lat', 'long', 'unix_time', 'merch_lat', 'merch_long'])
X = pd.get_dummies(X, columns=['merchant', 'category', 'gender', 'street', 'city'])
```

```
In [38]: from sklearn.preprocessing import MinMaxScaler
num_colss = ['amt', 'city_pop', 'Hour', 'Day_of_week', 'Month']
scaler = MinMaxScaler()
X[num_colss] = scaler.fit_transform(X[num_colss])
```

```
In [39]: y = X["is_fraud"]
X = X.drop("is_fraud", axis=1)
```

## Splitting into test and train data

```
In [40]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.4, random_state=42)
```

## Decision Tree Classifier

```
In [54]: from sklearn.tree import DecisionTreeClassifier
dt = DecisionTreeClassifier()
```

```
In [55]: dt.fit(X_train, y_train)
```

```
Out[55]: DecisionTreeClassifier()
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [56]: from sklearn.metrics import accuracy_score, mean_absolute_error ,mean_squared_error

print("Score of X-train with Y-train is : ", dt.score(X_train,y_train))
print("Score of X-test with Y-test is : ", dt.score(X_test,y_test))

y_pred=dt.predict(X_test)

print("Accuracy score " , accuracy_score(y_test,y_pred))

print("F1 score: ", round(f1_score(y_test, y_pred, average='weighted')*100,2))

print('Decision Tree:')
print('Accuracy:', accuracy_score(y_test, y_pred))
print('Classification Report:')
print(classification_report(y_test, y_pred))
print('Confusion Matrix:')
print(confusion_matrix(y_test, y_pred))
print('ROC AUC Score:')
y_prob = dt.predict_proba(X_test)[: , 1]
print(roc_auc_score(y_test, y_prob))
```

```
Score of X-train with Y-train is :  1.0
Score of X-test with Y-test is :  0.9978894376735375
Accuracy score  0.9978894376735375
F1 score:  99.78 %
Decision Tree:
Accuracy: 0.9978894376735375
Classification Report:
              precision    recall  f1-score   support

     0           1.00       1.00       1.00     230360
     1           0.74       0.67       0.70         858

 accuracy          0.99
 macro avg         0.87       0.83       0.85     231218
weighted avg         1.00       1.00       1.00     231218

Confusion Matrix:
[[230157    203]
 [   285    573]]
ROC AUC Score:
0.8334754692260337
```

## Logistic Regression

In [44]:

```

logreg = LogisticRegression()
logreg.fit(X_train, y_train)
y_pred = logreg.predict(X_test)
print('Logistic Regression:')
print('Accuracy:', accuracy_score(y_test, y_pred))
print('Classification Report:')
print(classification_report(y_test, y_pred))
print('Confusion Matrix:')
print(confusion_matrix(y_test, y_pred))
print('ROC AUC Score:')
y_prob = logreg.predict_proba(X_test)[:, 1]
print(roc_auc_score(y_test, y_prob))

```

C:\ProgramData\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:444: ConvergenceWarning: lbfgs failed to converge (status=2):  
ABNORMAL\_TERMINATION\_IN\_LNSRCH.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))

```
n_iter_i = _check_optimize_result(
```

Logistic Regression:

Accuracy: 0.9962892162374901

Classification Report:

C:\ProgramData\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1334: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

C:\ProgramData\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1334: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

C:\ProgramData\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1334: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	230360
1	0.00	0.00	0.00	858
accuracy			1.00	231218
macro avg	0.50	0.50	0.50	231218
weighted avg	0.99	1.00	0.99	231218

Confusion Matrix:

```
[[230360    0]
 [   858    0]]
```

ROC AUC Score:

0.5

```
In [47]: # Random Forest
rf = RandomForestClassifier(random_state=42)
rf.fit(X_train, y_train)
y_pred = rf.predict(X_test)
print('Random Forest:')
print('Accuracy:', accuracy_score(y_test, y_pred))
print('Classification Report:')
print(classification_report(y_test, y_pred))
print('Confusion Matrix:')
print(confusion_matrix(y_test, y_pred))
print('ROC AUC Score:')
y_prob = rf.predict_proba(X_test)[: , 1]
print(roc_auc_score(y_test, y_prob))
```

Random Forest:

Accuracy: 0.998209481960747

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	230360
1	0.98	0.53	0.69	858
accuracy			1.00	231218
macro avg	0.99	0.76	0.84	231218
weighted avg	1.00	1.00	1.00	231218

Confusion Matrix:

```
[[230350    10]
 [   404   454]]
```

ROC AUC Score:

0.9826292919039056